

10 Minutes to pandas

https://pandas.pydata.org/docs/user_guide/10min.html#object-creation

```
In [1]: import numpy as np
import pandas as pd
```

Object creation

```
In [2]: s = pd.Series([1,3,5,np.nan,6,8])
```

```
In [3]: s
```

```
Out[3]: 0    1.0
1    3.0
2    5.0
3    NaN
4    6.0
5    8.0
dtype: float64
```

```
In [4]: dates = pd.date_range("20130101", periods=6)
```

```
In [5]: dates
```

```
Out[5]: DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
                        '2013-01-05', '2013-01-06'],
                        dtype='datetime64[ns]', freq='D')
```

```
In [7]: df = pd.DataFrame(np.random.randn(6,4), index = dates, columns= list("ABCD"))
```

```
In [8]: df
```

```
Out[8]:
```

	A	B	C	D
2013-01-01	-1.398001	1.172644	-0.733226	1.460308
2013-01-02	-0.551749	-2.140217	0.965542	-0.458114
2013-01-03	-1.573659	-0.855627	0.560207	-0.702713
2013-01-04	-0.688304	-2.212931	-0.436499	0.724449
2013-01-05	0.203647	-1.176680	0.859715	0.762530
2013-01-06	-0.677257	-0.992430	-2.394600	-0.949607

```
In [9]: df2 = pd.DataFrame({
    "A": 1.0,
    "B": pd.Timestamp("20130102"),
    "C": pd.Series(1, index=list(range(4)), dtype="float32"),
    "D": np.array([3]*4, dtype="int32"),
```

```
"E":pd.Categorical(["test","train","test","train"]),
"F":"foo"
})
df2
```

```
Out[9]:
```

	A	B	C	D	E	F
0	1.0	2013-01-02	1.0	3	test	foo
1	1.0	2013-01-02	1.0	3	train	foo
2	1.0	2013-01-02	1.0	3	test	foo
3	1.0	2013-01-02	1.0	3	train	foo

```
In [10]: df2.dtypes
```

```
Out[10]: A      float64
B  datetime64[ns]
C      float32
D       int32
E      category
F      object
dtype: object
```

Viewing data

```
In [ ]: df.head()
```

```
In [ ]: df.tail(3)
```

```
In [ ]: df.index
# 행 축 제목 보기
```

```
In [ ]: df.columns
# 열 축 제목 보기
# 컬럼 보기
```

```
In [ ]: df2.to_numpy()
# 출력 시 행 열의 라벨을 포함하지 않는다.
```

```
In [ ]: df.describe()
```

```
In [ ]: df.T
# 행 열 전환
```

```
In [ ]: df.sort_index(axis=1, ascending = True)
```

```
In [ ]: df.sort_values(by="B")
```

Sslection

Getting

```
In [ ]: df["A"]  
# Series 호출
```

```
In [ ]: df[0:3]  
# 행 슬라이스
```

```
In [ ]: df["20130102":"20130104"]
```

Selection by Label

```
In [ ]: df.loc[dates[0]]
```

```
In [ ]: df.loc[:,["A","B"]]  
# 다중 축 호출
```

```
In [ ]: df.loc["20130102":"20130104",["A","B"]]
```

```
In [ ]: df.loc["20130102",["A","B"]]
```

```
In [ ]: df.loc[dates[0], "A"]
```

```
In [ ]: df.at[dates[0], "A"]
```

Selection by Position

```
In [ ]: df.iloc[3]
```

```
In [ ]: df.iloc[2]
```

```
In [ ]: df.iloc[3:5,0:2]
```

```
In [ ]: df.iloc[[1,2,4],[0,2]]
```

```
In [ ]: df.iloc[1:3,:]
```

```
In [ ]: df.iloc[:,1:3]
```

```
In [ ]:
```

```
df.iloc[1,1]
```

```
In [ ]: df.iat[1,1]
```

Boolean indexing

```
In [ ]: df[df["A"]>0]
# 해당 컬럼에서 조건을 만족하는 행만 추출
```

```
In [ ]: df[df>0]
```

```
In [ ]: df2 = df.copy()
```

```
In [ ]: df2["E"] = ["one", "one", "two", "three", "four", "three"]
df2
```

```
In [ ]: df2[df2["E"].isin(["two", "four"])]
# isin()
```

Setting

```
In [11]: s1 = pd.Series([1,2,3,4,5,6], index=pd.date_range("20130102", periods=6 ))
```

```
In [12]: s1
```

```
Out[12]: 2013-01-02    1
2013-01-03    2
2013-01-04    3
2013-01-05    4
2013-01-06    5
2013-01-07    6
Freq: D, dtype: int64
```

```
In [13]: df["F"] = s1
```

```
In [14]: df.at[dates[0], "A"] = 0
```

```
In [15]: df.iat[0,1] = 0
```

```
In [16]: df.loc[:, "D"] = np.array([5] * len(df))
df
```

```
Out[16]:
```

	A	B	C	D	F
2013-01-01	0.000000	0.000000	-0.733226	5	NaN
2013-01-02	-0.551749	-2.140217	0.965542	5	1.0

	A	B	C	D	F
2013-01-03	-1.573659	-0.855627	0.560207	5	2.0
2013-01-04	-0.688304	-2.212931	-0.436499	5	3.0
2013-01-05	0.203647	-1.176680	0.859715	5	4.0
2013-01-06	-0.677257	-0.992430	-2.394600	5	5.0

```
In [17]: df2 = df.copy()
```

```
In [18]: df2[df2>0] = -df2
df2
```

Out [18]:

	A	B	C	D	F
2013-01-01	0.000000	0.000000	-0.733226	-5	NaN
2013-01-02	-0.551749	-2.140217	-0.965542	-5	-1.0
2013-01-03	-1.573659	-0.855627	-0.560207	-5	-2.0
2013-01-04	-0.688304	-2.212931	-0.436499	-5	-3.0
2013-01-05	-0.203647	-1.176680	-0.859715	-5	-4.0
2013-01-06	-0.677257	-0.992430	-2.394600	-5	-5.0

Missing data

```
In [19]: df1 = df.reindex(index=dates[0:4], columns=list(df.columns) + ["E"])
df1.loc[dates[0] : dates[1], "E"] = 1
df1
```

Out [19]:

	A	B	C	D	F	E
2013-01-01	0.000000	0.000000	-0.733226	5	NaN	1.0
2013-01-02	-0.551749	-2.140217	0.965542	5	1.0	1.0
2013-01-03	-1.573659	-0.855627	0.560207	5	2.0	NaN
2013-01-04	-0.688304	-2.212931	-0.436499	5	3.0	NaN

```
In [20]: df1.dropna(how="any")
# 결측치 있는 행 제거
```

Out [20]:

	A	B	C	D	F	E
2013-01-02	-0.551749	-2.140217	0.965542	5	1.0	1.0

```
In [21]: df1.fillna(value= 5)
# 결측치 채우기
```

Out [21]:

	A	B	C	D	F	E
--	---	---	---	---	---	---

	A	B	C	D	F	E
2013-01-01	0.000000	0.000000	-0.733226	5	5.0	1.0
2013-01-02	-0.551749	-2.140217	0.965542	5	1.0	1.0
2013-01-03	-1.573659	-0.855627	0.560207	5	2.0	5.0
2013-01-04	-0.688304	-2.212931	-0.436499	5	3.0	5.0

```
In [22]: pd.isna(df1)
# 결측치 확인
```

```
Out[22]:
```

	A	B	C	D	F	E
2013-01-01	False	False	False	False	True	False
2013-01-02	False	False	False	False	False	False
2013-01-03	False	False	False	False	False	True
2013-01-04	False	False	False	False	False	True

Operation

Stats

```
In [23]: df.mean()
```

```
Out[23]: A    -0.547887
B    -1.229648
C    -0.196477
D     5.000000
F     3.000000
dtype: float64
```

```
In [24]: df.mean(1)
```

```
Out[24]: 2013-01-01    1.066694
2013-01-02    0.854715
2013-01-03    1.026184
2013-01-04    0.932453
2013-01-05    1.777336
2013-01-06    1.187142
Freq: D, dtype: float64
```

```
In [25]: s = pd.Series([1,3,5,np.nan, 6,8],index=dates).shift(2)
```

```
In [26]: s
```

```
Out[26]: 2013-01-01    NaN
2013-01-02    NaN
2013-01-03     1.0
2013-01-04     3.0
2013-01-05     5.0
2013-01-06    NaN
Freq: D, dtype: float64
```

```
In [27]: df.sub(s, axis="index")
```

```
Out[27]:
```

	A	B	C	D	F
2013-01-01	NaN	NaN	NaN	NaN	NaN
2013-01-02	NaN	NaN	NaN	NaN	NaN
2013-01-03	-2.573659	-1.855627	-0.439793	4.0	1.0
2013-01-04	-3.688304	-5.212931	-3.436499	2.0	0.0
2013-01-05	-4.796353	-6.176680	-4.140285	0.0	-1.0
2013-01-06	NaN	NaN	NaN	NaN	NaN

Apply

```
In [29]: df.apply(np.cumsum)
```

```
Out[29]:
```

	A	B	C	D	F
2013-01-01	0.000000	0.000000	-0.733226	5	NaN
2013-01-02	-0.551749	-2.140217	0.232316	10	1.0
2013-01-03	-2.125407	-2.995844	0.792523	15	3.0
2013-01-04	-2.813712	-5.208775	0.356024	20	6.0
2013-01-05	-2.610065	-6.385455	1.215739	25	10.0
2013-01-06	-3.287322	-7.377886	-1.178861	30	15.0

```
In [32]: df.apply(lambda x: x.max() - x.min())
```

```
Out[32]: A    1.777306
         B    2.212931
         C    3.360142
         D    0.000000
         F    4.000000
         dtype: float64
```

Histogramming

https://pandas.pydata.org/docs/user_guide/basics.html#basics-discretization

```
In [33]: s = pd.Series(np.random.randint(0, 7, size=10))
         s
```

```
Out[33]: 0    3
         1    2
         2    2
         3    3
         4    5
         5    6
         6    0
         7    1
         8    2
         9    0
         dtype: int32
```

```
In [34]: s.value_counts()
```

```
Out[34]: 2    3
         0    2
         3    2
         1    1
         5    1
         6    1
         dtype: int64
```

String Methods

```
In [35]: s = pd.Series(["A", "B", "C", "Aaba", "Baca", np.nan, "CABA", "dog", "Cat"])
         s.str.lower()
```

```
Out[35]: 0      a
         1      b
         2      c
         3    aaba
         4    baca
         5     NaN
         6    caba
         7    dog
         8    cat
         dtype: object
```

Merge

https://pandas.pydata.org/docs/user_guide/merging.html#merging

Concat

```
In [37]: df = pd.DataFrame(np.random.randn(10,4))
         df
```

```
Out[37]:
```

	0	1	2	3
0	0.112079	-0.909524	-0.082731	0.370589
1	0.825064	0.418057	0.524716	0.913962
2	1.671239	1.200347	1.125526	0.819722
3	0.511291	0.611885	1.797537	-0.805969
4	1.031593	-0.531125	1.705441	1.013867
5	1.031899	-0.162714	-0.817580	0.865577
6	0.125774	-0.531272	-1.615048	-0.081109
7	0.139489	-0.124460	-0.498288	0.607304
8	-0.639346	-0.858044	2.055711	-0.325547
9	-0.057988	-3.002323	-0.104391	0.762115

```
In [38]: pieces = [df[:3], df[3:7], df[7:]]
         pd.concat(pieces)
```

```
Out[38]:
```


	0	1	2	3
0	0.112079	-0.909524	-0.082731	0.370589
1	0.825064	0.418057	0.524716	0.913962
2	1.671239	1.200347	1.125526	0.819722
3	0.511291	0.611885	1.797537	-0.805969
4	1.031593	-0.531125	1.705441	1.013867
5	1.031899	-0.162714	-0.817580	0.865577
6	0.125774	-0.531272	-1.615048	-0.081109
7	0.139489	-0.124460	-0.498288	0.607304
8	-0.639346	-0.858044	2.055711	-0.325547
9	-0.057988	-3.002323	-0.104391	0.762115

- 데이터 프레임의 컬럼을 합치는 건 상대적으로 빠르다.
- 하지만 행을 합치는 건 리소스 소모가 심하다.

Join

```
In [40]: left = pd.DataFrame({"key":["foo","foo"], "lval":[1,2]})
right = pd.DataFrame({"key":["foo","foo"], "rval":[4,5]})
left
```

```
Out[40]:
```

	key	lval
0	foo	1
1	foo	2

```
In [41]: right
```

```
Out[41]:
```

	key	rval
0	foo	4
1	foo	5

```
In [42]: pd.merge(left, right, on = "key")
```

```
Out[42]:
```

	key	lval	rval
0	foo	1	4
1	foo	1	5
2	foo	2	4
3	foo	2	5

```
In [43]: left = pd.DataFrame({"key":["foo","bar"], "lval":[1,2]})
right = pd.DataFrame({"key":["foo","bar"], "rval":[4,5]})
```

```
left
```

```
Out[43]:
```

	key	lval
0	foo	1
1	bar	2

```
In [44]:
```

```
right
```

```
Out[44]:
```

	key	rval
0	foo	4
1	bar	5

```
In [45]:
```

```
pd.merge(left, right, on="key")
```

```
Out[45]:
```

	key	lval	rval
0	foo	1	4
1	bar	2	5

Grouping

https://pandas.pydata.org/docs/user_guide/groupby.html#groupby

```
In [46]:
```

```
df = pd.DataFrame(
    {
        "A": ["foo", "bar", "foo", "bar", "foo", "bar", "foo", "foo"],
        "B": ["one", "one", "two", "three", "two", "two", "one", "three"],
        "C": np.random.randn(8),
        "D": np.random.randn(8),
    }
)
df
```

```
Out[46]:
```

	A	B	C	D
0	foo	one	-0.529658	0.498554
1	bar	one	1.965949	0.338279
2	foo	two	-0.211798	-1.796187
3	bar	three	-0.032970	-2.755438
4	foo	two	-1.850905	-0.245048
5	bar	two	0.034017	-0.539988
6	foo	one	-1.089231	2.067523
7	foo	three	1.549236	1.837700

```
In [47]:
```

```
df.groupby("A").sum()
# 범주형 데이터는 결과물에서 제외 되었음
```

Out[47]:

	C	D
A		
bar	1.966996	-2.957147
foo	-2.132355	2.362541

In [48]:

```
df.groupby(["A", "B"]).sum()
```

Out[48]:

		C	D
A	B		
bar	one	1.965949	0.338279
	three	-0.032970	-2.755438
	two	0.034017	-0.539988
foo	one	-1.618889	2.566077
	three	1.549236	1.837700
	two	-2.062703	-2.041235

Reshaping

https://pandas.pydata.org/docs/user_guide/advanced.html#advanced-hierarchical

https://pandas.pydata.org/docs/user_guide/reshaping.html#reshaping-stacking

Stack

In [51]:

```
tuples = list(
    zip(
        *[
            ["bar", "bar", "baz", "baz", "foo", "foo", "qux", "qux"],
            ["one", "two", "one", "two", "one", "two", "one", "two"],
        ]
    )
)

index = pd.MultiIndex.from_tuples(tuples, names=["first", "second"])
df = pd.DataFrame(np.random.randn(8,2), index=index, columns=["A", "B"])
df2 = df[:4]
df2
```

Out[51]:

		A	B
first	second		
bar	one	1.083845	-0.705348
	two	0.822413	-0.066388
baz	one	0.029070	0.293777
	two	1.398663	-0.374066

```
In [52]: stacked = df2.stack()
```

```
In [53]: stacked
```

```
Out[53]: first second
bar one A 1.083845
      one B -0.705348
      two A 0.822413
      two B -0.066388
baz one A 0.029070
      one B 0.293777
      two A 1.398663
      two B -0.374066
dtype: float64
```

```
In [54]: stacked.unstack()
```

```
Out[54]:
```

		A	B
first	second		
bar	one	1.083845	-0.705348
	two	0.822413	-0.066388
baz	one	0.029070	0.293777
	two	1.398663	-0.374066

```
In [55]: stacked.unstack(1)
```

```
Out[55]:
```

	second	one	two
first			
bar	A	1.083845	0.822413
	B	-0.705348	-0.066388
baz	A	0.029070	1.398663
	B	0.293777	-0.374066

```
In [56]: stacked.unstack(0)
```

```
Out[56]:
```

	first	bar	baz
second			
one	A	1.083845	0.029070
	B	-0.705348	0.293777
two	A	0.822413	1.398663
	B	-0.066388	-0.374066

Pivot tables

```
In [58]: df = pd.DataFrame(
```

```

.....: {
.....:     "A": ["one", "one", "two", "three"] * 3,
.....:     "B": ["A", "B", "C"] * 4,
.....:     "C": ["foo", "foo", "foo", "bar", "bar", "bar"] * 2,
.....:     "D": np.random.randn(12),
.....:     "E": np.random.randn(12),
.....: }
.....: )
df

```

Out[58]:

	A	B	C	D	E
0	one	A	foo	-0.378944	0.508452
1	one	B	foo	-0.574639	-0.006769
2	two	C	foo	-1.290909	0.028426
3	three	A	bar	0.412040	-1.094314
4	one	B	bar	0.044262	-0.001531
5	one	C	bar	2.971559	1.464310
6	two	A	foo	-0.011837	0.082054
7	three	B	foo	-0.096568	1.146086
8	one	C	foo	-0.209786	2.609733
9	one	A	bar	1.464482	0.027279
10	two	B	bar	1.198594	0.555783
11	three	C	bar	-1.505907	-1.343206

In [59]:

```
pd.pivot_table(df, values="D", index=["A", "B"], columns=["C"])
```

Out[59]:

		C	bar	foo
A	B			
one	A	1.464482	-0.378944	
	B	0.044262	-0.574639	
	C	2.971559	-0.209786	
three	A	0.412040		NaN
	B		NaN	-0.096568
	C	-1.505907		NaN
two	A		NaN	-0.011837
	B	1.198594		NaN
	C		NaN	-1.290909

Time series

In [61]:

```

rng = pd.date_range("1/1/2012", periods=100, freq="S")
ts = pd.Series(np.random.randint(0,500, len(rng)), index=rng)
ts.resample("5Min").sum()

```

```
Out[61]: 2012-01-01    26519
Freq: 5T, dtype: int32
```

```
In [62]: rng = pd.date_range("3/6/2012 00:00", periods=5, freq="D")
ts = pd.Series(np.random.randn(len(rng)), rng)
ts
```

```
Out[62]: 2012-03-06    -0.882651
2012-03-07     0.033077
2012-03-08    -0.594378
2012-03-09    -0.438064
2012-03-10     1.221211
Freq: D, dtype: float64
```

```
In [63]: ts_utc = ts.tz_localize("UTC")
ts_utc
```

```
Out[63]: 2012-03-06 00:00:00+00:00    -0.882651
2012-03-07 00:00:00+00:00     0.033077
2012-03-08 00:00:00+00:00    -0.594378
2012-03-09 00:00:00+00:00    -0.438064
2012-03-10 00:00:00+00:00     1.221211
Freq: D, dtype: float64
```

```
In [64]: ts_utc.tz_convert("US/Eastern")
```

```
Out[64]: 2012-03-05 19:00:00-05:00    -0.882651
2012-03-06 19:00:00-05:00     0.033077
2012-03-07 19:00:00-05:00    -0.594378
2012-03-08 19:00:00-05:00    -0.438064
2012-03-09 19:00:00-05:00     1.221211
Freq: D, dtype: float64
```

```
In [65]: rng = pd.date_range("1/1/2012", periods = 5, freq="M")
ts = pd.Series(np.random.randn(len(rng)), index = rng)
ts
```

```
Out[65]: 2012-01-31     0.191414
2012-02-29     0.796353
2012-03-31    -0.173852
2012-04-30     0.274467
2012-05-31    -2.236095
Freq: M, dtype: float64
```

```
In [66]: ps = ts.to_period()
ps
```

```
Out[66]: 2012-01     0.191414
2012-02     0.796353
2012-03    -0.173852
2012-04     0.274467
2012-05    -2.236095
Freq: M, dtype: float64
```

```
In [67]: ps.to_timestamp()
```

```
Out[67]: 2012-01-01     0.191414
2012-02-01     0.796353
2012-03-01    -0.173852
2012-04-01     0.274467
```

2012-05-01 -2.236095
 Freq: MS, dtype: float64

```
In [68]: prng = pd.period_range("1990Q1", "2000Q4", freq = "Q-NOV")
         ts = pd.Series(np.random.randn(len(prng)), prng)
         ts.index = (prng.asfreq("M", "e") + 1).asfreq("H", "s") + 9
         ts.head()
```

```
Out[68]: 1990-03-01 09:00    1.463594
         1990-06-01 09:00   -0.521600
         1990-09-01 09:00    0.057676
         1990-12-01 09:00   -0.095379
         1991-03-01 09:00    0.424294
         Freq: H, dtype: float64
```

Categoricals

```
In [69]: df = pd.DataFrame({
         "id": [1,2,3,4,5,6] , "raw_grade" : ["a","b","b","a","a","e"]
       })

         df["grade"] = df["raw_grade"].astype("category")
         df["grade"]

         ## 범주화
```

```
Out[69]: 0    a
         1    b
         2    b
         3    a
         4    a
         5    e
         Name: grade, dtype: category
         Categories (3, object): ['a', 'b', 'e']
```

```
In [70]: df["grade"].cat.categories = ["Very Good", "Good", "Very Bad"]
         df["grade"] = df["grade"].cat.set_categories(
             ["Very Bad", "Bad", "Medium", "Good", "Very Good"])

         df["grade"]
```

```
Out[70]: 0    Very Good
         1         Good
         2         Good
         3    Very Good
         4    Very Good
         5    Very Bad
         Name: grade, dtype: category
         Categories (5, object): ['Very Bad', 'Bad', 'Medium', 'Good', 'Very Good']
```

```
In [72]: df.sort_values(by="grade")
```

```
Out[72]:
```

	id	raw_grade	grade
5	6	e	Very Bad
1	2	b	Good
2	3	b	Good
0	1	a	Very Good
3	4	a	Very Good

	id	raw_grade	grade
4	5	a	Very Good

```
In [73]: df.groupby("grade").size()
```


```
Out[73]: grade
Very Bad    1
Bad         0
Medium      0
Good        2
Very Good   3
dtype: int64
```

Ploting

```
In [74]: import matplotlib.pyplot as plt
plt.close("all")
```


```
In [75]: ts = pd.Series(np.random.randn(1000), index=pd.date_range("1/1/2000", periods = 1000))
ts = ts.cumsum()
ts.plot()
```

```
Out[75]: <AxesSubplot:>
```




```
In [76]: df = pd.DataFrame(
    np.random.randn(1000,4), index = ts.index, columns=["A","B","C","D"])
df = df.cumsum()
plt.figure()
df.plot()
plt.legend(loc="best")
```

```
Out[76]: <matplotlib.legend.Legend at 0x1b7240d74f0>
<Figure size 432x288 with 0 Axes>
```

```
In [79]: plt.figure()  
df.plot()  
plt.legend(loc='best')
```

```
Out[79]: <matplotlib.legend.Legend at 0x1b724305370>  
<Figure size 432x288 with 0 Axes>
```



Getting data in/out

CSV

```
In [ ]: df.to_csv("fadfa.csv")
```